

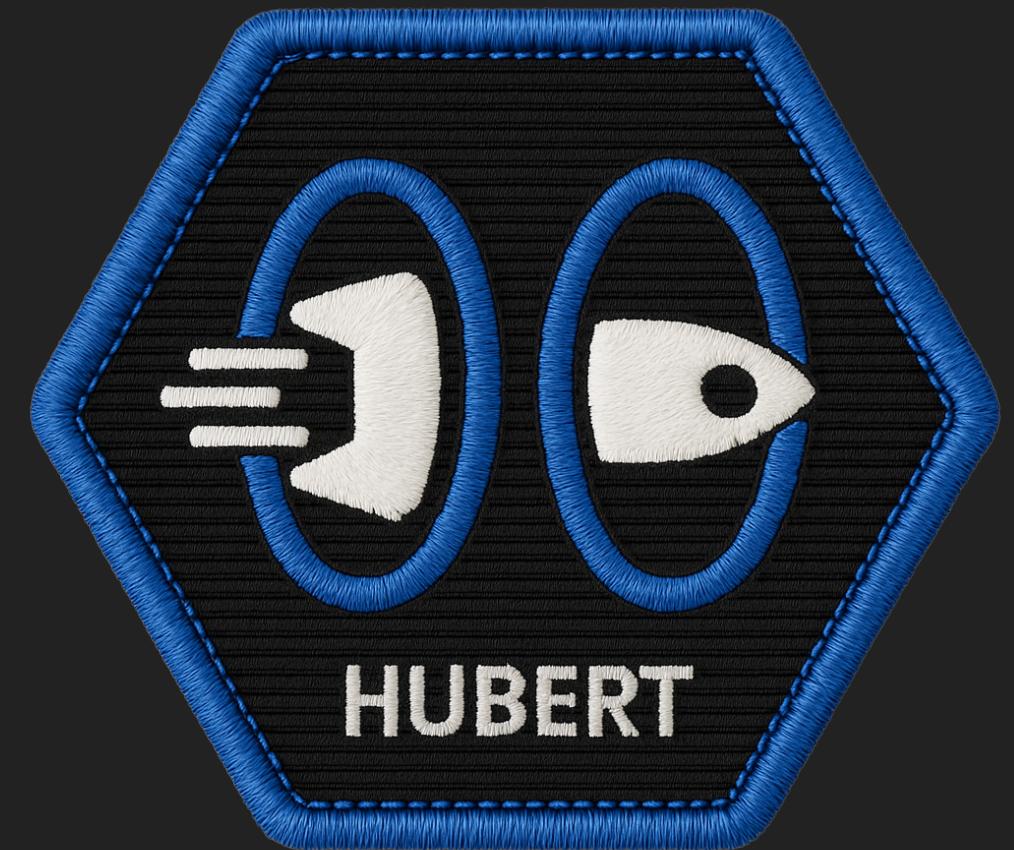
BLOCKCHAIN COMMONS

HUBERT



WHAT IS HUBERT?

- ▶ A protocol that facilitates *secure multiparty transactions*:
- ▶  **Participants write once** using random keys
- ▶  **Messages contain random keys** for expected responses, enabling bidirectional communication
- ▶  **Complete opacity** to outsiders through end-to-end encryption
- ▶  **No central server** required for coordination
- ▶  **Trustless operation** using public distributed networks



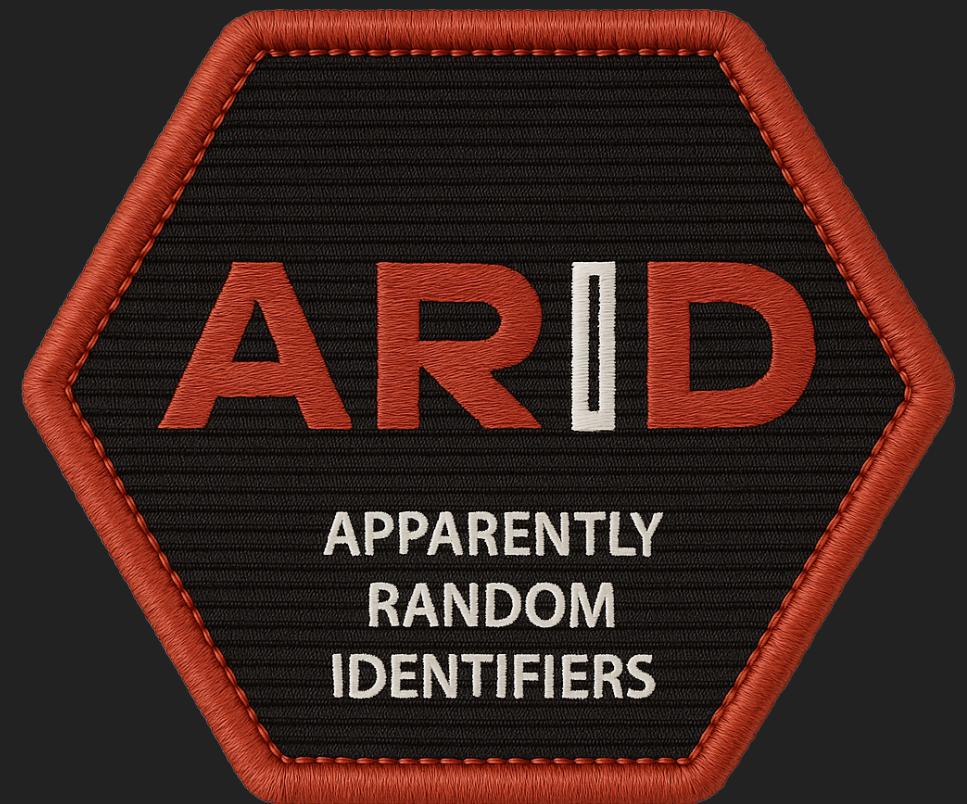
WHAT ARE URS?

- ▶ UR: Uniform Resource
 - ▶ Defined in **BCR-2020-005**
 - ▶ <https://github.com/blockchaincommons/research>
 - ▶ Encodes binary data as typed, easy to handle text URI
 - ▶ `ur:type/bytewords`



WHAT IS AN ARID?

- ▶  **ARID**: Apparently Random Identifier
- ▶ Defined in **BCR-2022-002**
 - ▶ <https://github.com/blockchaincommons/research>
 - ▶ 256 statistically random bits (32 bytes)
 - ▶ Can refer to anything
 - ▶ But cannot be correlated to anything
 - ▶ In Hubert, ARIDs are addresses of cryptographic *dead drops*.



WHAT IS A DEAD DROP?

- ▶ Spycraft: A way to exchange messages where the parties never have to meet in person.
- ▶ Sender leaves message at known location, receiver visits later and picks it up.
- ▶ Requires out-of-band coordination to set up locations, protocols, and signals.



WHAT IS GORDIAN ENVELOPE?

- ▶ “Smart document” format for structured data with a built-in Merkle-like digest tree.
- ▶ Core superpower: hashed-data elision (selective redaction) with verifiable inclusion proofs—disclose only what’s needed while signatures remain valid.
- ▶ Radically recursive: envelopes contain other envelopes.
- ▶ Privacy and human-rights oriented design (RFC 6973 / RFC 8280 alignment).



WHAT IS GORDIAN ENVELOPE?

- ▶ Flexible cryptography—Schnorr, Ed25519, SSH, Post-Quantum algorithms supported for signing/encryption.
- ▶ Supports non-correlation via salt.
- ▶ Suited to cryptographic seeds, keys, verifiable credentials, and other high-integrity artifacts
- ▶ Compact binary format with minimal overhead.
- ▶ Wire-agnostic: integrity, privacy, and semantics travel with the data—ideal for offline/air-gapped and multi-hop workflows.



WHAT IS GORDIAN SEALED TRANSACTION PROTOCOL (GSTP)?

- ▶ Secure, transport-agnostic request/response protocol built on Gordian Envelope
- ▶ Works over HTTP/TCP, Bluetooth, NFC, QR “sneakernet,” Tor, etc.
- ▶ Messages are both encrypted to recipients and signed by the sender
- ▶ Provides Encrypted State Continuations (ESC) to carry encrypted workflow state inside the messages themselves.



WHAT IS GORDIAN SEALED TRANSACTION PROTOCOL (GSTP)?

- ▶ Asynchronous and peer-to-peer friendly; also supports classic client-server flows and multimodal (multi-channel) transactions.
- ▶ Encourages idempotent actions so retries over flaky links are safe.
- ▶ Designed to be “wire-agnostic”: encryption/signing are at the protocol layer, so insecure/unreliable transports are acceptable.



STORAGE LAYERS



BitTorrent™



IPFS



STORAGE LAYERS



BitTorrent™



STORAGE: BITTORRENT MAINLINE DHT



- ▶ Used by BitTorrent clients to find peers.
- ▶ In 2013, the concurrent number of users of Mainline DHT was 16-28 million, with intra-day changes of at least 10 million.
- ▶ UDP overlay where peers collectively replace trackers; keys live in a 160-bit space and lookups walk “closest” nodes.
- ▶ “Sloppy” DHT—eventually consistent and best-effort. Items expire without re-announcement.
- ▶ Signaling/coordination layer, not durable storage.



STORAGE: BITTORRENT MAINLINE DHT



- ▶ Two data modes:
 - ▶ Immutable items (key = SHA-1 of the value)
 - ▶ Mutable items (key = SHA-1 of the publisher's Ed25519 public key concatenated with optional salt)
- ▶ Values are practically capped at ~1 KB.
- ▶ Expect messages to last ~30-45 minutes without republishing



STORAGE: BITTORRENT MAINLINE DHT



- ▶ We use the “mutable” mode by stretching the ARID into an Ed25519 key pair.
- ▶ We derive the Ed25519 private key from the ARID deterministically using HKDF:
 - ▶ Private key used for signing the payload
 - ▶ Public key derived from private key, used to compute DHT storage location and verify payload
 - ▶ Write-once semantics enforced by rejecting updates past seq=1



STORAGE: IPFS



- ▶ IPFS (InterPlanetary File System) is a distributed, content-addressed storage network
- ▶ Established network with consistent operation; well-suited for coordination and signaling. ~56,000 dedicated infrastructure nodes globally
- ▶ Content persists reliably when pinned; 48-hour DHT expiration ensures fresh routing data
- ▶ Content addressing: Files identified by CID (Content Identifier) = cryptographic hash of content



STORAGE LAYERS



IPFS



STORAGE: IPFS



- ▶ Two-layer storage model:
 - ▶ Immutable layer: Data stored at CID (content hash), retrieved via Kademlia DHT
 - ▶ Mutable layer: IPNS (InterPlanetary Name System) provides mutable pointers to CIDs
- ▶ Size limits: Practical limit ~1-10 MB per item (much larger than DHT's 1 KB)
- ▶ Persistence: Content expires from DHT in 48 hours without re-announcement
 - ▶ IPNS records should be republished every 4 hours (Kubo default)
 - ▶ Long-term persistence requires "pinning" (explicit storage on specific nodes)
- ▶ Latency: 1-10 seconds for IPNS resolution (slower than Mainline DHT)



STORAGE: IPFS



- ▶ Dependencies: Requires Kubo daemon or compatible IPFS node with RPC API
- ▶ Write-once semantics: Before publishing, check if IPNS name resolves to any CID
 - ▶ If resolution succeeds → name already published → reject with AlreadyExists error
 - ▶ If resolution fails with “could not resolve name” → name unpublished → proceed with publish
 - ▶ Unlike Mainline DHT’s sequence number approach, relies on IPNS name existence check



STORAGE LAYERS



BitTorrent™



IPFS



STORAGE: HYBRID



- ▶ Hybrid storage combines Mainline DHT and IPFS with automatic size-based routing
- ▶ Intelligent storage selection:
 - ▶ Small envelopes (≤ 1000 bytes): Stored directly in Mainline DHT
 - ▶ Large envelopes (> 1000 bytes): Two-step indirection process
- ▶ Storage process for large envelopes:
 - ▶ 1. Generate new "reference ARID" for the actual envelope
 - ▶ 2. Store actual envelope in IPFS at reference ARID
 - ▶ 3. Store small reference envelope in DHT at original ARID



STORAGE: HYBRID



- ▶ Retrieval process:
 - ▶ 1. Fetch envelope from DHT using original ARID
 - ▶ 2. Check if it's a reference envelope (has 'dereferenceVia': "ipfs")
 - ▶ 3. If reference: Extract reference ARID and fetch actual envelope from IPFS
 - ▶ 4. If not reference: Return DHT envelope directly



STORAGE: HYBRID



- ▶ Benefits:
 - ▶ Transparent to caller: Application uses same API regardless of storage backend
 - ▶ Fast retrieval for small messages (DHT: 1-5 seconds)
 - ▶ Large capacity for big messages (IPFS: up to 10 MB)
 - ▶ No size calculation required by caller



STORAGE: HYBRID



- ▶ Dependencies: Requires both embedded DHT client and Kubo daemon for IPFS
- ▶ Write-once semantics: Enforced separately by each backend
 - ▶ (DHT seq=1, IPFS existence check)
- ▶ Still to be implemented:
 - ▶ Encryption of reference envelope to multiple recipients; currently stored as plaintext and exposes IPFS ARID.



STORAGE LAYERS



STORAGE: SERVER



- ▶ Centralized HTTP-based storage for testing and controlled deployments
- ▶ Architecture: Simple HTTP POST API with write-once semantics
 - ▶ PUT endpoint: /put (stores ARID + envelope + optional TTL)
 - ▶ GET endpoint: /get (retrieves envelope by ARID with polling)
- ▶ Two storage backends:
 - ▶ Memory: In-memory HashMap, volatile (lost on restart)
 - ▶ SQLite: Persistent file-based database storage



STORAGE: SERVER



- ▶ Storage features:
 - ▶ Direct ARID-to-envelope mapping (no key derivation like DHT/IPFS)
 - ▶ Write-once enforcement: Returns 409 CONFLICT if ARID already exists
 - ▶ TTL support: Optional time-to-live in seconds for automatic expiration
 - ▶ Background cleanup: Expired entries pruned automatically (every 60 seconds for SQLite)





DEMO

HUBERT